



Language Modelling with N-grams

MMAI 5400 – lecture 3
Fall 2024



Contents

Why is NLP difficult?

- Challenges with natural language

Intro to language models

- N-gram language models
 - Limitations
 - Fixes
 - Evaluation



Book chapter 3 – N-gram Language Models

Read

- Sections: 3 - 3.7 & 3.9

Don't read (although, I'm not going to stop you)

- Section: 3.8 *Advanced: Perplexity's Relation to Entropy*

NLP challenges – part 1

Context

- The meaning of natural language is often context dependent

Ambiguity (multiple meanings)

- Crash blossoms ([crash blossom - Wiktionary, the free dictionary](#)), examples:


“Violinist Linked to JAL Crash Blossoms”

“Red Tape Holds Up New Bridges”

“Local High School Dropouts Cut in Half”

Natural languages → ≈infinite number of possible sentences

- Listing all sentence–meaning pairs is not possible



NLP challenges – part 2



Non-standard language

“U taught us 2 @neversaynever!”

“Were SOO PROUD of what youve accomplished!”

Ambiguous segmentation

The New York-New Haven railroad vs The New York-New Haven railroad

NLP challenges – part 3



Idioms

“Get cold feet”

“lose face”

“raining cats & dogs”

Neologisms

“Unfriend”

“Retweet”

“bromance”

Difficult entity names

*“Where is **A Bug’s Life** playing?”*

*“when was **Let it be** recorded?”*

*“... a mutation in the **for** gene...”*



Two ways of using text in machine learning

Bag-of-Words

- Based on word (token) counts
- Documents are represented by their word frequencies
- Word order is ignored
- Good for long documents

Sequence models

- Mainly neural network-based
- Documents are represented as a sequence of tokens (word/sub-word/character)
- Word order is important
- Works sentence level



Two ways of using text in machine learning

Bag-of-Words

- Based on word (token) counts
- Documents are represented by their word frequencies
- Word order is ignored
- Good for long documents

Class 4

Sequence models

- Mainly neural network-based
- Documents are represented as a sequence of tokens (word/sub-word/character)
- Word order is important
- Works sentence level

Classes 8 to 12

Language models & N-grams



Language models (LMs)

Models that assign probabilities to sequences of words (e.g. sentences)

$$P(w_1, w_2, \dots, w_t)$$

Simple LM

- N-grams

Neural LM

- Autoregressively trained RNNs & Transformers



Language Models: linguistic prophets

Which is the more natural-sounding sentence?

Whoever is happy will make others happy too.

Whoever make is will happy happy too others.



Language Models: linguistic prophets

Question: What word would you use to complete the following sentence?

What have you been up to ____

- *beehive*
- *lately*
- *N-gram*
- *complete*



Language Models: linguistic prophets

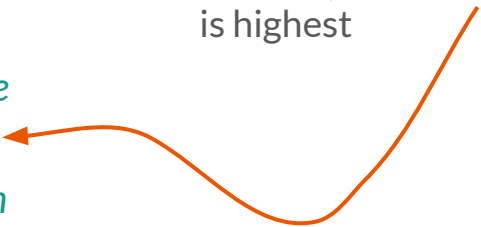
Question: What word would you use to complete the following sentence?

What have you been up to ____

- *What have you been up to **beehive***
- *What have you been up to **lately***
- *What have you been up to **N-gram***
- *What have you been up to **complete***

Most likely
I.e. among the options

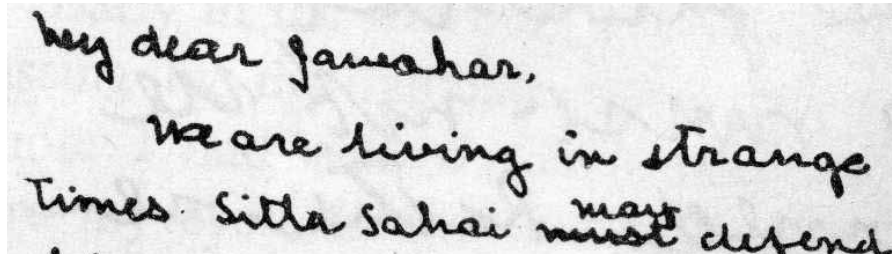
$P(\text{what, have, you, been, up, to, lately})$
is highest



Why are LMs useful?

Question: What do these two examples have in common?

"Their are two finals."



My dear Jawahar,
We are living in strange
times. Sitla Sahai ^{may} must defend

Image credit: Mohandas K. Gandhi - <http://www.mkgandhi.org/images/lefthand.JPG>, Public Domain,
<https://commons.wikimedia.org/w/index.php?curid=1455138>

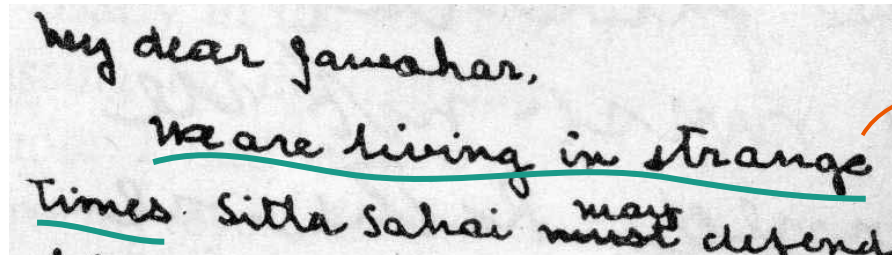
Why are LMs useful?

Question: What do these two examples have in common?

- We use our knowledge of likely sentences to correct & disambiguate

"Their are two finals."

"There are two finals."



A photograph of a handwritten note on aged paper. The text is written in cursive and reads: "My dear Jawahar, we are living in strange times. Sitla Sahai may must defend". The phrase "we are living in strange times" is underlined in blue ink. The word "may" is written above "must".

"We are living in strange times"

When are LMs useful?

- Examples

Human computer interfacing (lecture 12)

Optical character recognition (OCR)

Spelling & grammar correction

Chatbots

Machine translation (MMAI 5500 lecture 9)

Augmentative & alternative communication systems

Text generation (lectures 11 & 12)





Simple Language Models (LM)

- N-grams

N-gram: a sequence of N words

1-gram examples

1-gram (unigram)

- “wake”, “up”

2-gram (bigram)

2-gram examples

3-gram (trigram)

- “wake up”, “this morning”, ...

...

3-gram examples

N-gram

- “wake up this”, “up this morning”

Question: Is the chain rule below related to the chain rule used to compute gradients?

N-gram LM

Probability of a sequence



Probability of an n word long sequence: $P(w_1 w_2 \dots w_n)$

Count

The probability of a specific n word long sequence can be computed by:

$$P(w_1 w_2 \dots w_n) = \frac{C(w_1 w_2 \dots w_n)}{C(\text{all } n \text{ word long sequences})}$$


However, it is hard to count the number of all n word long sequences.

It is easier to use *the chain-rule of probability*:


$$P(w_1 w_2 \dots w_n) = P(w_1)P(w_2|w_1)P(w_3|w_1, w_2)\dots P(w_n|w_1 \dots w_{n-1})$$

N-gram LM




The chain-rule of probability:

A product of conditional probabilities


$$P(w_1 w_2 \dots w_n) = P(w_1)P(w_2|w_1)P(w_3|w_1 w_2)\dots P(w_n|w_1 \dots w_{n-1})$$

Joint probability written more concisely:


$$P(w_1 \dots w_n) = P(w_{1:n})$$

A sequence from 1 to n

The chain-rule written more concisely:

$$P(w_{1:n}) = \prod_{t=1}^n P(w_t|w_{1:t-1})$$

N-gram example



Bigram: $P(w_t|w_{t-1})$

example:

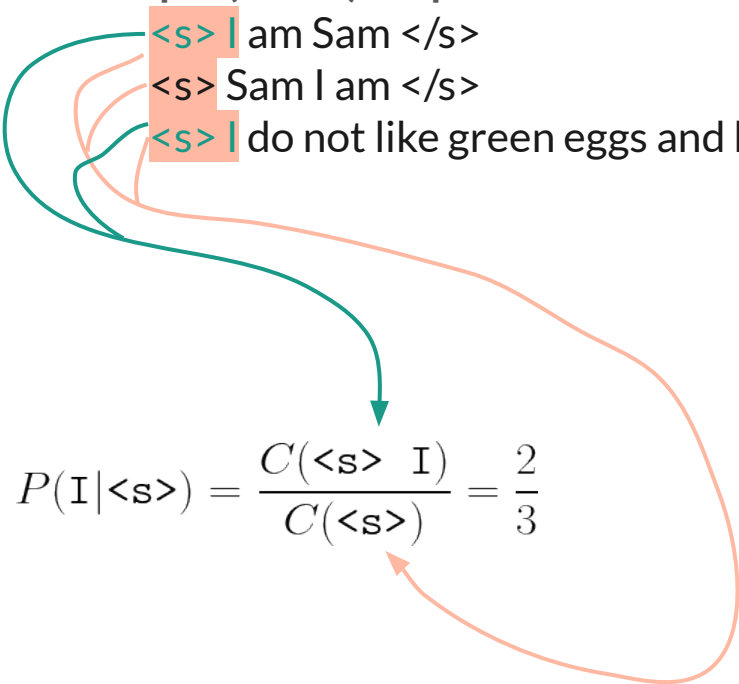
$\langle s \rangle$ is the start token
 $\langle /s \rangle$ is the end token

Example (mirco) corpus:

$\langle s \rangle$ I am Sam $\langle /s \rangle$

$\langle s \rangle$ Sam I am $\langle /s \rangle$

$\langle s \rangle$ I do not like green eggs and ham $\langle /s \rangle$



N-gram example



Example (mirco) corpus:

<s> I am Sam </s>

<s> Sam I am </s>

<s> I do not like green eggs and ham </s>

Bigram: $P(w_t|w_{t-1})$

examples:

$$P(I|<s>) = \frac{C(<s> I)}{C(<s>)} = \frac{2}{3}$$

$$P(do|I) = \frac{C(I do)}{C(I)} = \frac{1}{3}$$

$$P(do|not) = \frac{C(do not)}{C(do)} = \frac{1}{1}$$

$$P(</s>|Sam) = \frac{C(Sam</s>)}{C(Sam)} = \frac{1}{2}$$

Trigram: $P(w_t|w_{t-1} w_{t-2})$

examples:

$$P(do|<s> I) = \frac{C(<s> I do)}{C(<s> I)} = \frac{1}{2}$$

$$P(am|Sam I) = \frac{C(Sam I am)}{C(Sam I)} = \frac{1}{1} = 1$$

The probability of the sequence is approximated with the N-gram probabilities.

The LM – $P(w_{1:n})$

Bigram LM

$$P(w_{1:n}) \approx \prod_{t=1}^n P(w_t | w_{t-1})$$

$$\begin{aligned} & P(< s> \text{ I do not}) \\ & \approx P(< s> \text{ I})P(\text{I do})P(\text{do not}) \\ & = \frac{C(< s> \text{ I})}{C(< s>)} \frac{C(\text{I do})}{C(\text{I})} \frac{C(\text{do not})}{C(\text{do})} \\ & = \frac{2}{3} \times \frac{1}{3} \times \frac{1}{1} \\ & \approx 0.22 \end{aligned}$$

Example (mirco) corpus:

<s> I am Sam </s>

<s> Sam I am </s>

<s> I do not like green eggs and ham </s>

Trigram LM

$$P(w_{1:n}) \approx \prod_{t=1}^n P(w_t | w_{t-2:t-1})$$

$$\begin{aligned} & P(< s> \text{ I am Sam } </s>) \\ & \approx P(< s> < s> \text{ I})P(< s> \text{ I am})P(\text{I am Sam})P(\text{am Sam } </s>) \\ & = \frac{C(< s> < s> \text{ I})}{C(< s> < s>)} \frac{C(< s> \text{ I am})}{C(< s> \text{ I})} \frac{C(\text{I am Sam})}{C(\text{I am})} \frac{C(\text{am Sam } </s>)}{C(\text{am Sam})} \\ & = \frac{2}{3} \times \frac{1}{2} \times \frac{1}{2} \times \frac{1}{1} \\ & \approx 0.17 \end{aligned}$$

Generating sentences with language models



Bigram example

1. Start with a *start token*: $\langle s \rangle$
2. Randomly draw a new word according to $P(w^i | \langle s \rangle) \forall i \in V$
3. Append the word w to the sequence
4. Randomly draw the next word according to $P(w_t^i | w_{t-1}) \forall i \in V$
5. Repeat steps 3 & 4
6. End when the *end token*, $\langle /s \rangle$, is drawn

Start ($\langle s \rangle$) and end ($\langle /s \rangle$) tokens are part of the vocabulary V

The vocabulary V is a list of unique/distinct words in the dataset.

Question: What is the most likely sequence, starting with <s>?

Generating a sequence

- example

Vocabulary

$$V = \{\text{go}, \text{</s>}, \text{here}, \text{we}, \text{<s>}\}$$

Bigram probabilities

$w_{t-1} \backslash w_t$	go	</s>	here	we	<s>
go	0.05	0	0.35	0.2	0.15
</s>	0.40	0	0.30	0.2	0.05
here	0.25	0	0.30	0.4	0.35
we	0.30	0	0.05	0.2	0.45
<s>	0	0	0	0	0



Limits of N-grams

Would a good **bi-gram** model assign the following “sentence” high or low probability?

It is it is it is it is it is it is it is.

Why?



Limits of N-grams

Would a good **tri-gram** model assign the following “sentence” high or low probability?

It is it is it is it is it is it is it is.

Why?



Limits of N-grams

- data requirement

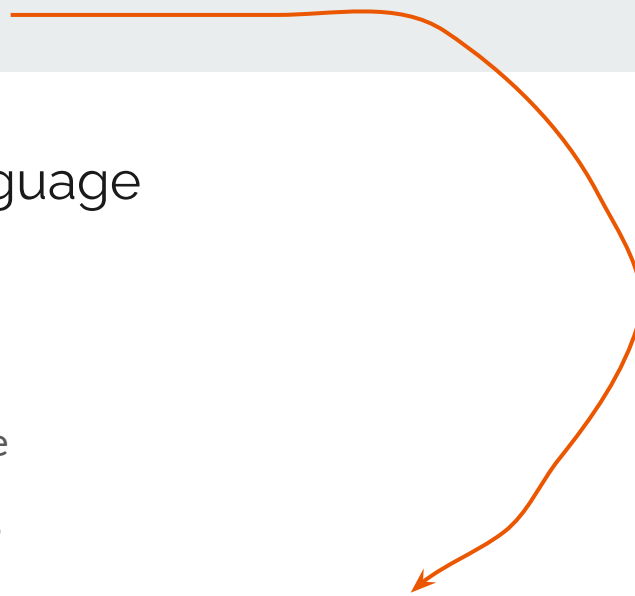
Higher order N-grams capture longer range dependencies between words

But, the number of unique possibilities of 3, 4, & 5-word sequences increases exponentially

- An example of the combinatorial explosion

Thus, **exponentially more data is required to get non-zero counts** of all/nearly all sequences

Question: Which of the sentences below sound most likely?



Limits of N-grams

- the generative property of language

Natural language is generative.

I.e. we can generate novel sentences that still make sense

Examples. Which of the sentences below are most likely?

- *“The arctic char is Lapland’s most beautiful fish.” vs “The beautiful most arctic Lapland’s fish char is.”*
- *“Is, angry Norway & the small suicidal adorable lemming.” vs “The suicidal Norway lemming is small, angry & adorable.”*

Limits of N-grams


- the generative property of language

Natural language is generative.

I.e. we can generate novel sentences that still make sense

Examples. Which of the sentences below are most likely?

- *“The arctic char is Lapland’s most beautiful fish.” vs “The beautiful most arctic Lapland’s fish char is.”*
- *“Is, angry Norway & the small suicidal adorable lemming.” vs “The suicidal Norway lemming is small, angry & adorable.”*



Limits of N-grams

- zero counts*

Training set (**bigram continuation**: count)

- **denied the** *allegations*: 5
- **denied the** *speculation*: 2
- **denied the** *rumors*: 1
- **denied the** *report*: 1

Hypothetical test set (**bigram continuation**)

- **denied the** *offer*
- **denied the** *loan*

Two problems: 1) obviously sensible 3-grams were assigned 0 probability; 2) *perplexity* is undefined

What to do with sequences (e.g. trigrams) missing from the training set?

Dealing with unknown words



Closed vocabulary

- The test set can only contain words from a fixed size, known vocabulary

Open vocabulary

- Potentially unknown words can occur in the test set
- We allow for unknown/out-of-vocabulary words
- All unknown words are replaced by the <UNK> token

Unknown word

- Words that occur in the test set, but **NOT** in the training set
- Or (not xor) words that are very rare in training set

Two strategies

- Replace all words not in a fixed vocabulary with <UNK>
- Replace all words below some frequency with <UNK>



Dealing with zero counts

- smoothing

Add-1 smoothing aka Laplace smoothing

- Add 1 to all N-gram counts & normalize by the number of N-gram counts
- Results in a smoother probability distribution with no zeroes & lower max
- With many zeroes, too much probability mass is shifted resulting in bad estimates
- Not good performance & is not in common use



Dealing with zero counts

- smoothing

Backoff

- If counts for a higher-order N-gram is missing, then use the counts for a lower-order instead
- E.g. if the count for the tri-gram w_{n-2}, w_{n-1}, w_n is missing, then we can use the count for w_{n-1}, w_n instead.
- Example: if the sequence “*suicidal Norway lemming*” does not occur then we can back off to “*Norway lemming*”

Further improvement: we can use a weighted average over all N-grams (uni, bi, tri, ...) in place of a high-order N-gram (called interpolation in the textbook)



Google N-grams

Released 2006 (before deep learning NLP)

1 to 5-grams

From 1,024,908,267,229 words of running text

Based on all 5-word sequences (1,176,470,663) that appear at least 40 times

For a data set of this size smoothing has to be fast, & not too sophisticated


Google N-gram viewer

Question: Why do you think "wake up" has a higher probability than "up this morning"?

Google Books Ngram Viewer

Graph these comma-separated phrases: ☐ case-insensitive

between and from the corpus with smoothing of [Search lots of books](#)



Language models

- Evaluation



Extrinsic

- Deploy the LM in an application & see how much the application improves
- Only way of knowing if the LM really helps the task
- E.g.
 - Task: speech recognition
 - Two LMs
 - Run the task with each model & compare results

Intrinsic

- Use an evaluation metric that is independent of the application
- Evaluate on a test set
- Best model: the LM assigning the highest probability to the test set



Actually: lowest perplexity

Language models

- Evaluation metric

Instead of raw probabilities we use **perplexity** (PP)

- PP is the inverse probability of the test set normalized by the number of words

$$\text{test set } W_{test} = w_1, w_2, \dots, w_N$$

$$PP(W_{test}) = \sqrt[N]{\frac{1}{P(w_1, w_2, \dots, w_N)}}$$

Language models

- Generalization



How to select training data?

- An LM will reflect the corpus (data set) it is trained on
- Trivial example: an LM trained on a corpus of *Swedish* texts will be useless for *English* sentence generation
- Better example: an LM trained on Shakespeare won't do good on tasks involving modern legal documents

Solutions

1. Train on a corpus similar to the future application
2. Train on a huge corpus containing all kinds of texts (e.g. GPT-3* was trained on [2 TB text](#))

* GPT-3 is a big *neural* LM aka an LLM



Terminology

Language model

Vocabulary

N-gram

- The N-gram vs the N-gram LM

Perplexity

Extrinsic evaluation

Intrinsic evaluation

Chain-rule of probability

LLM



Book exercises

Do the following exercises from the book:

- 3.1
- 3.2
- 3.3
- 3.4

You will be tested on similar problems on the midterm



Tutorial

`MMAI5400_class03_ngrams.ipynb`