



Linguistic features

MMAI 5400 – lecture 6
Fall 2024



Contents

Sentence-level tasks

- Uses of parsing

Part-of-speech tagging

Dependency parsing

- Parse trees
- Shallow parsing
- Dependency structure

Why do we need dependency parsing?

- Examples 😊

Evaluation

Tutorial



NLP tasks

There are many different NLP tasks.

[NLP-progress: Tracking Progress in Natural Language Processing](#) lists the tasks, descriptions & state-of-the-art (SOTA) results.



A few sentence-level tasks

POS tagging

- Tag words with their part-of-speech (POS)

Coreference resolution

- Clustering mentions in text that refer to the same underlying real-world entities.
- Example: *I* voted for Trudeau because he was most aligned with my values", she said.
- http://nlpprogress.com/english/coreference_resolution.html



A few sentence-level tasks

Dependency parsing

- Describe the syntactic structure of a sentence in terms of the words (or lemmas) in a sentence & associated directed binary grammatical relations among the words.
- Approximates semantic relationships thus making it useful for e.g. *coreference resolution*, *question answering* & *information extraction*.
- http://nlpprogress.com/english/dependency_parsing.html

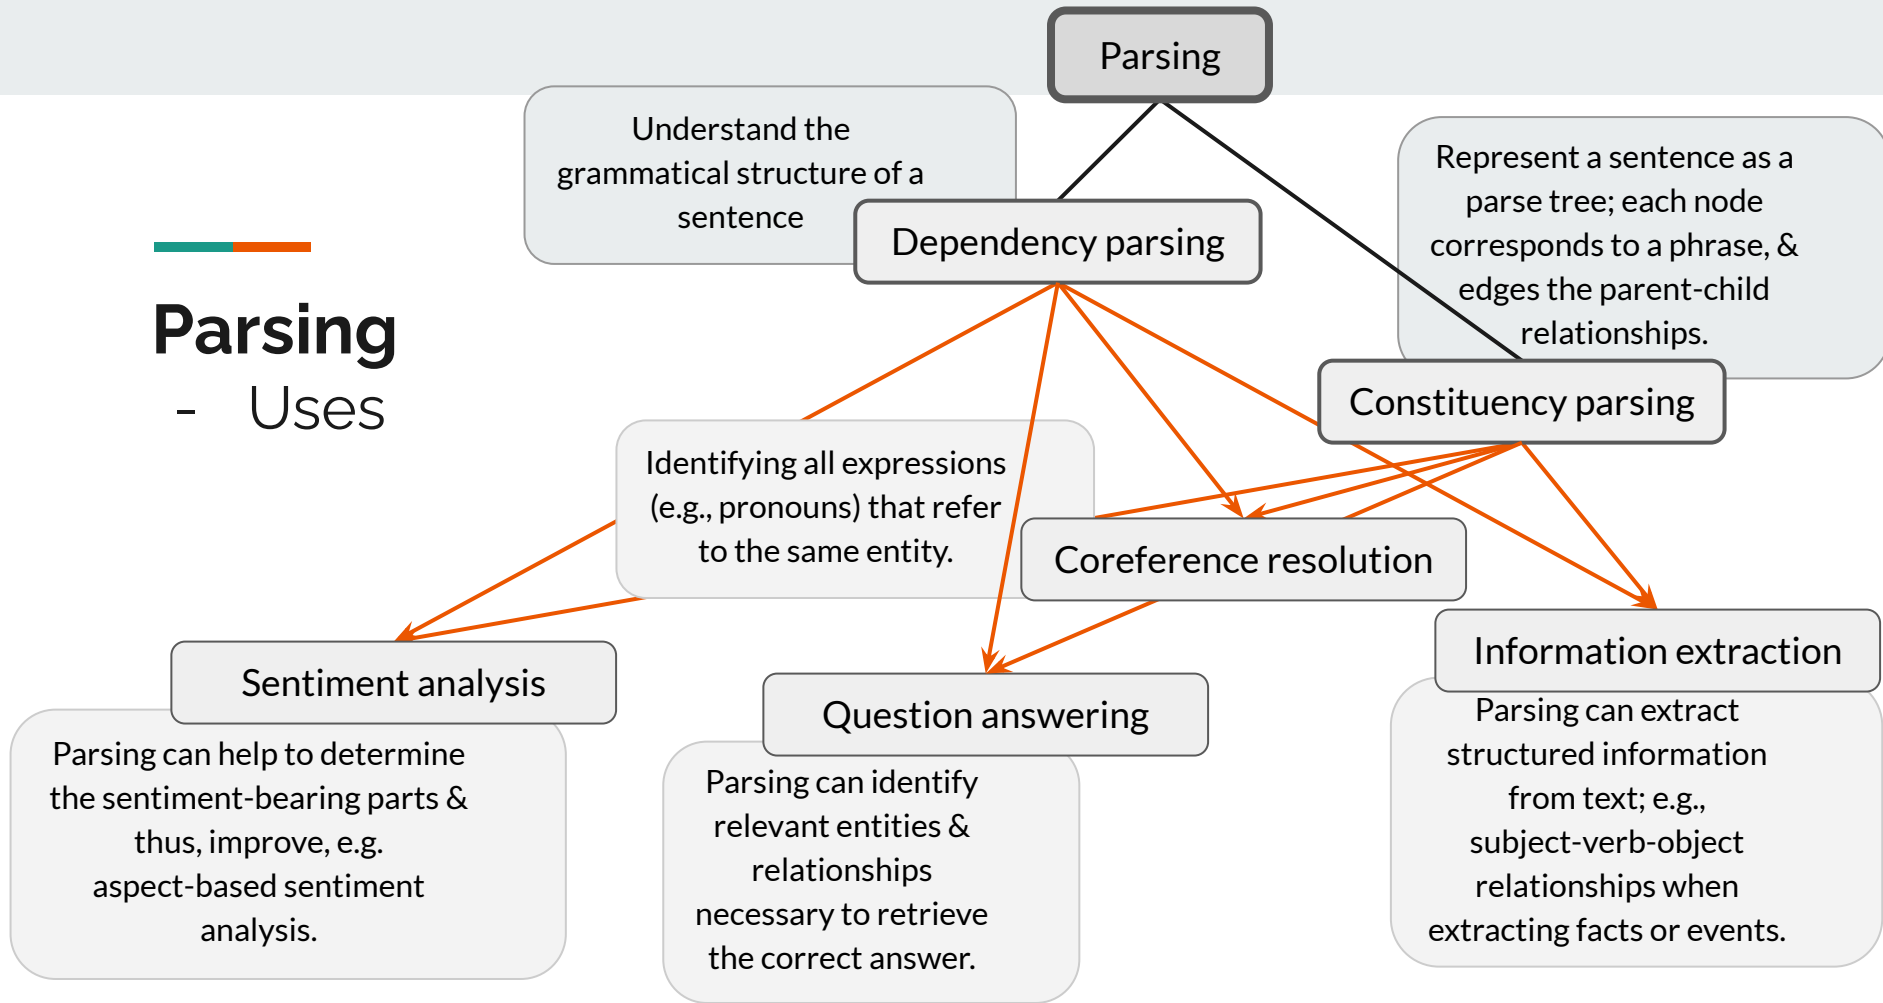
No need to know this

NSUBJ	-	Nominal subject
DOBJ	-	Direct object
DET	-	Determinant
NMOD	-	Nominal modifier
CASE	-	Prepositions, postpositions & other case markers



Parsing

- Uses



Part-of-speech Tagging



POS tagging aka grammatical tagging


Part-of-speech (POS) aka word class, or syntactic category.

A POS is a category of words with similar grammatical properties. Common English parts of speech are *noun, verb, adjective, adverb, pronoun, preposition, conjunction*, etc.

POS tagging is a disambiguation task since many words have more than one possible POS.

POS is an important step in many applications. Examples are *lemmatization, Named Entity Recognition (NER), sentiment analysis, question answering, word sense disambiguation*, & for grammar & spell checkers.

POS tagging



NP - proper noun, singular

NN - noun, singular

VBD - verb, past tense

DT - determinant

JJ - adjective



Ambiguity

- 6 different POS for the word *back*

Examples of ambiguous words:
that, back, down, put & set.

earnings growth took a **JJ** **back** seat

a small building in the **NN** **back**

a clear majority of senators **VBD** **back** the bill

Dave began to **VB** **back** toward the door

enable the country to buy **RP** **back** about debt

I was twenty-one **RB** **back** then

The Penn Treebank POS tagset

A total of 45 tags

Only a subset in the tables below, see the complete list in Figure 8.2 of *Speech & Language Processing*

CC	coordinating conjunction	<i>and</i>
CD	cardinal number	<i>1, third</i>
DT	determiner	<i>the</i>
EX	existential there	<i>there is</i>
FW	foreign word	<i>cachaça</i>
JJ	adjective	<i>green</i>
NNP	proper noun, singular	<i>Hjalmar</i>
NPS	proper noun, plural	<i>Vikings</i>

NN	Noun, singular or mass	<i>llama</i>
VB	verb be, base form	<i>be</i>
VBD	verb be, past tense	<i>was, were</i>
VV	verb, base form	<i>take</i>
VVD	verb, past tense	<i>took</i>
RB	adverb	<i>quickly</i>
RP	particle	<i>up, off</i>
MD	modal	<i>could, will, would</i>



Two PoS tagging algorithms

For simplicity, we will only use 3 simplified tags:

- **N** noun: e.g. *house, bike, rock, llama, Hjalmar*
 - Notice that we have fused several noun-related tags
- **M** modal verb: e.g. *will, could, would, should, may*
- **V** verb: e.g. *see, run, jump, skip, ride*
 - Notice that we have fused several verb-related tags



The baseline algorithm: most frequent tag

Corpus:

N **V** **N**
Liv saw Adam
N **V** **N**
Adam saw Will

Tag counts

	N	V
Liv	1	0
Adam	2	0
Will	1	0
saw	0	2



The baseline algorithm: most frequent tag

Corpus:

N **V** **N**
Liv saw Adam
N **V** **N**
Adam saw Will

New sentence:

Liv saw Will

Tag counts

	N	V
Liv	1	0
Adam	2	0
Will	1	0
saw	0	2



The baseline algorithm: most frequent tag

Corpus:

N **V** **N**
Liv saw Adam
N **V** **N**
Adam saw Will

New sentence:

N **V** **N**
Liv saw Will

Tag counts

	N	V
Liv	1	0
Adam	2	0
Will	1	0
saw	0	2



The baseline algorithm: most frequent tag

Corpus:

N **V** **N**
Liv saw Adam
N **V** **N**
Adam saw Will

New sentence:

N **V** **N**
Liv saw Will

Tag counts

	N	V
Liv	1	0
Adam	2	0
Will	1	0
saw	0	2

This is the baseline algorithm that any new algorithm has to improve upon

- The most-frequent-tag baseline achieves 92.34% accuracy on WSJ, whereas SOTA is around 97%

Question: What is the problem?



A problem with the baseline algorithm

Corpus:

N M V N

Liv will see Adam.

N M V N

Will will see Adam.

N M V N

Adam will see Liv.


New sentence:

Liv will see Will.

Tag counts

	N	M	V
Liv	2	0	0
Adam	3	0	0
Will	1	3	0
see	0	0	3

Question: What context can we use to disambiguate between the tags?



A problem with the baseline algorithm

Document

N M V N

Liv will see Adam.

N M V N

Will will see Adam.

N M V N

Adam will see Liv.

New sentence:

N M V M


Liv will see Will.

Tag counts

	N	M	V
Liv	2	0	0
Adam	3	0	0
Will	1	3	0
see	0	0	3

Will is associated with 2 different tags,
i.e. it is ambiguous - **context is needed**

Question: What context can we use to disambiguate between the tags?



Solving this ambiguity is for better tagging algorithms, but not for this lecture.

A problem with the baseline algorithm

Document

N M V N

Liv will see Adam.

N M V N

Will will see Adam.

N M V N

Adam will see Liv.

New sentence:

N M V M

Liv will see Will.

Tag counts

	N	M	V
Liv	2	0	0
Adam	3	0	0
Will	1	3	0
see	0	0	3

Will is associated with 2 different tags, i.e. it is ambiguous - **context is needed**



POS-tagged corpora

The Brown corpus

- A million words taken from 500 texts of different genres.

The WSJ corpus

- A million words taken from Wall Street Journal.

The Switchboard corpus

- 2 million words from telephone conversations.

Tagging was done by an automatic tagger, followed by human corrections.

Dependency Parsing



Wen dependency parsing?

Mainly for improving the performance of downstream tasks, not an application in it self.

With end-to-end models it can be used to enrich the text representation (e.g. marking up text before vectorizing, or as an additional feature vector).

A good parsing model is trained on very large datasets. Thus, a custom model fine-tuned on a small dataset can leverage this important knowledge. Think about it as transferring knowledge from large datasets to small-dataset models (i.e. transfer learning).



Language is recursive

Starting units: words with POS categories.

the,	cat,	cuddly,	door,	by
Det	N	Adj	N	P

Words combine into phrases with categories.

the	cuddly	cat,	by	the	door
Det	Adj	N	P	Det	N
		NP	P		NP

Phrases can combine into bigger phrases recursively.

the	cuddly	cat	by	the	door
		NP	P		NP
					NP

Parse trees

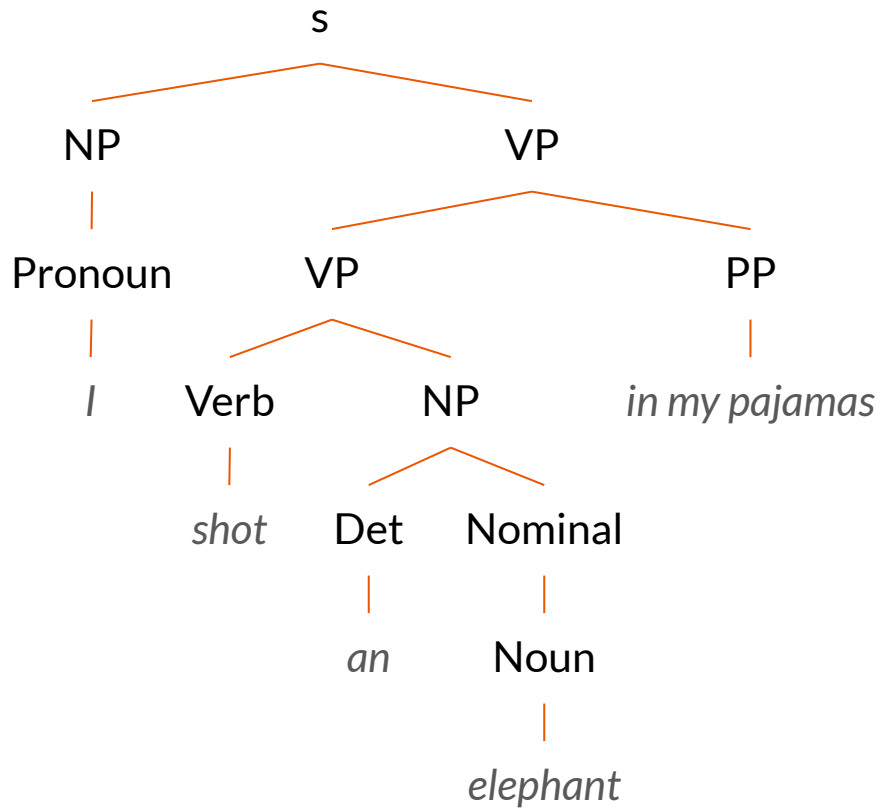
NP - *Noun Phrase*, a sequence of words surrounding a noun.

VP - *Verb Phrase*, a verb & other constituents, e.g. NP & PP.

Det - *Determiner* (aka article), e.g. **a** stop, **the** flights, etc.

Nominal - follows the determiner & contains pre- & post-noun modifiers. Can consist of only a single noun.

PP - *Prepositional Phrase*, a preposition followed by an NP.



Shallow parsing



Aka **chunking**, chunks groups of adjacent tokens into phrases on the basis of their POS tags.
Well-known chunks are *noun phrases*, *verb phrases*, & *prepositional phrases*.

Noun phrase

- A group of at least 2 words with a noun (e.g. 'she', 'he', 'they', ...) or a pronoun as its head & includes modifiers (e.g. 'the', 'a', 'of them', 'with her').
- A noun phrase plays the role of a noun.
- Examples (modifier **noun**):
 - the **man**
 - a **girl**
 - the **doggy** in the window

Verb phrase

- A group of at least 2 words consisting of a verb (the head) & arguments that further illustrates the verb tense, action or tone.
- Examples (arguments **verb**):
 - She can **smell** the pizza
 - He has **appeared** on screen as an actor.



Dependency structure

Dependency structure shows which words *depend* on which other words.


- *Depend*: modify or are arguments of.

Look in the large crate in the kitchen by the door

Dependency structure

Dependency structure shows which words *depend* on which other words.

- *Depend*: modify or are arguments of.



Dependency structure

Dependency structure shows which words *depend* on which other words.

- *Depend*: modify or are arguments of.



Look for the large barking dog by the door in a crate

**Why do we need
dependency parsing?**

—



Why is it important to understand the structure of sentences?

Sentence structure is important to be able to correctly interpret language.

Humans communicate by composing words together into bigger units capable of conveying complex meanings.

Thus, for many NLP tasks it becomes crucial to know what is connected to what, & in what way.

- I.e. what words are arguments & modifiers of other words.

subj - subject
obj - object
nmod - nominal modifier

How it can go wrong

- Prepositional phrase attachment ambiguity

2 alternative meanings.

- The man who was killed had a knife.
- Cops lethally stab a man.



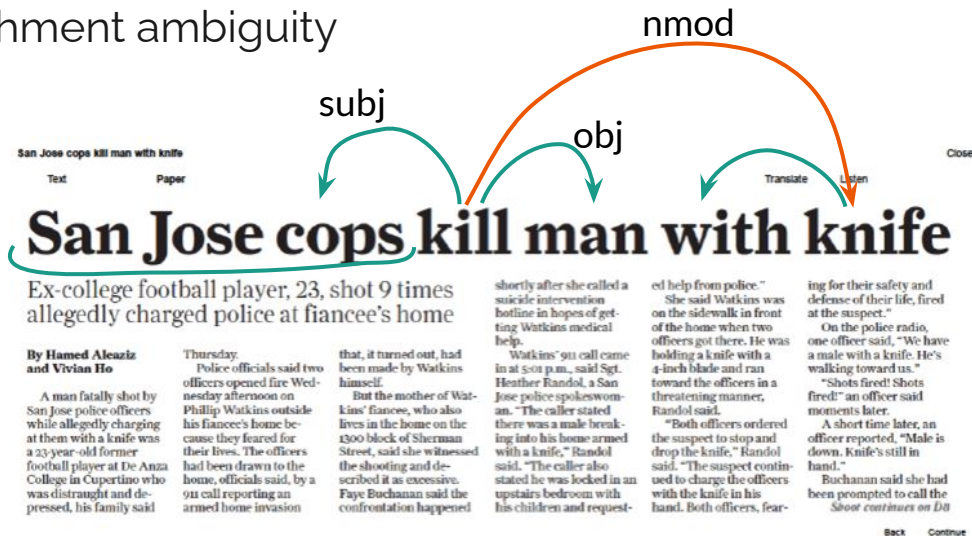
subj - subject
obj - object
nmod - nominal modifier

How it can go wrong

- Prepositional phrase attachment ambiguity

2 alternative meanings.

- The man who was killed had a knife.
- Cops lethally stab a man.



Question: What are the 2 interpretations of this headline?

How it can go wrong

- Prepositional phrase attachment ambiguity



The screenshot shows the BBC News website interface. At the top, there is a navigation bar with the BBC logo, a 'Sign in' button, and links to Home, News, Sport, Reel, Worklife, and Travel. Below this is a red banner with the word 'NEWS' in white. Underneath the banner is a secondary navigation bar with links to Home, Coronavirus, Video, World, US & Canada, UK, Business, Tech, Science, Stories, and Entertainment & Arts. The 'Science' link is highlighted with a red underline. The main headline is 'Scientists count whales from space' in a large, bold, dark font. Below the headline, it says 'By Jonathan Amos' and 'BBC Science Correspondent'. At the bottom, it shows a clock icon and the date '1 November 2018'.

BBC Sign in Home News Sport Reel Worklife Travel

NEWS

Home | Coronavirus | Video | World | US & Canada | UK | Business | Tech | Science | Stories | Entertainment & Arts

Science

Scientists count whales from space

By Jonathan Amos
BBC Science Correspondent

🕒 1 November 2018

How it can go wrong

- Prepositional phrase attachment ambiguity






Image credits:

<https://www.nrdc.org/onearth/space-whales-look-much-smaller>

<https://www.laphamsquarterly.org/roundtable/fantastical-allure-space-whale>

Example from the
Wall Street Journal



How it can go wrong

- Coordination scope ambiguity

Shuttle veteran & longtime NASA executive Fred Gregory appointed to board.

Shuttle veteran & longtime NASA executive Fred Gregory appointed to board.

Question: What are the 2 interpretations of this headline?

Question: What is the cause of the ambiguity?

How it can go wrong

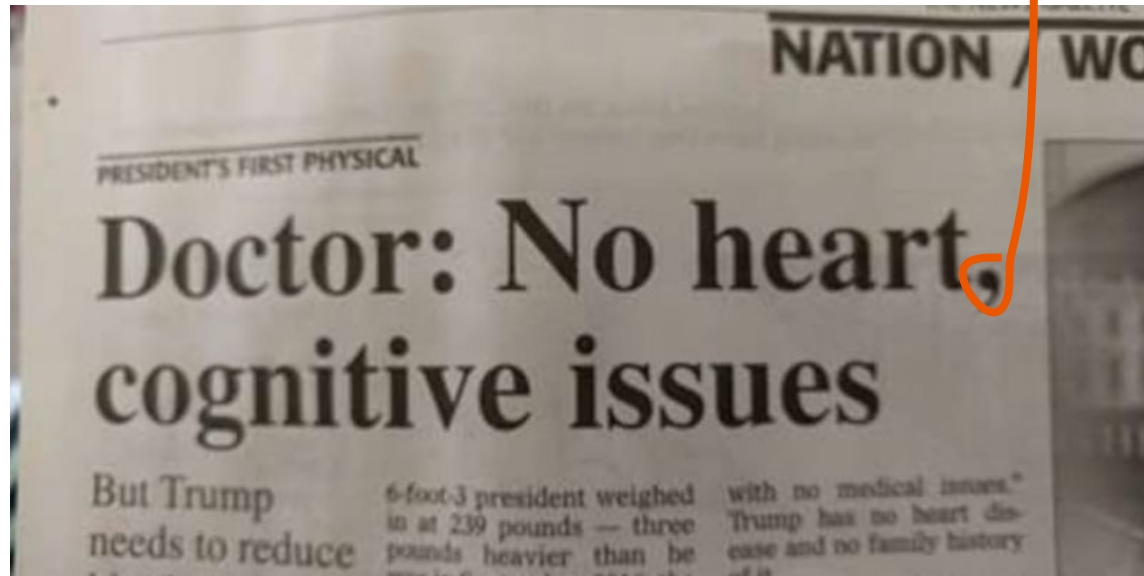
- Coordination scope ambiguity



How it can go wrong

- Coordination scope ambiguity

&/or



Question: What are the 2 interpretations of this headline?

Question: What is the cause of the ambiguity?

When it goes wrong

- Verb Phrase (VP) attachment ambiguity



When it goes wrong

- Verb Phrase (VP) attachment ambiguity



Dependency grammar



Dependency grammar

Dependency grammar assumes that sentence structure consists of relations between words, normally binary asymmetric relations (i.e. arrows) called *dependencies*.

The arrows are often *typed* with the name of grammatical relations (e.g. subject, determiner, prepositional object, etc).

Long history

The first linguist Pāṇini (around 6 - 4th century BCE) creates something like the first dependency grammar.

Written on birch bark!!

Earliest NLP dependency parser is from 1962 (Hays, 1962).



Image credit: CC-BY-4.0, CC BY 4.0

<https://commons.wikimedia.org/w/index.php?curid=36055831>

<https://wellcomeimages.org/indexplus/image/L0032691.html>

https://wellcomeimages.org/indexplus/obf_images/f8/3a/5863acae2cbb1fc99ef08cc8f13e.jpg

Grammatical relations in dependency parsing

Universal dependency relations


See [Universal Dependencies](#) & [Stanford typed dependencies manual](#) for a complete list.

Causal argument relations	Description
NSUBJ	Nominal subject
DOBJ	Direct object
IOBJ	Indirect object
CCOMP	Clausal complement
XCOMP	Open clausal complement

Nominal modifier relations	Description
NMOD	Nominal modifier
AMOD	Adjectival modifier
NUMMOD	Numeric modifier
APPOS	Appositional modifier
DET	Determiner
CASE	Prepositions, postpositions & other case markers
Other notable relations	Description
CONJ	Conjunction
CC	Coordinating conjunction

Dependency parsing

- A sentence is parsed by deciding for each word what other word (including a fake ROOT) it is a dependent of.
- Constraints are often added
 - Only one word is a dependent of ROOT.
 - No cycles, i.e. $A \rightarrow B$ & $B \rightarrow A$.
- The above constraints makes the dependencies a tree (a directed acyclic graph).
- Sometimes arrows are not allowed to cross (non-projective).





Methods for dependency parsing

Dynamic programming

- Eisner (1996) gives a clever algorithm with complexity $O(n^3)$, by producing parse items with heads at the ends rather than in the middle.

Graph algorithms

- Create a Minimum Spanning Tree (MST) for a sentence.
- McDonald et al.'s (2005) MSTParser scores dependencies independently using an ML classifier.

Constraint Satisfaction

- Edges are eliminated that don't satisfy hard constraints. Karlsson (1990), etc.


Transition-based parsing

- Greedy choice of dependencies guided by ML classifiers, e.g. MaltParser (Nivre et al. 2008). Popular & effective.

Evaluating a dependency parser



Evaluation – UAS



Unlabeled Attachment Score (UAS)

Gold (i.e. ground truth)

to	from	word	type
1	2	She	nsubj
2	0	saw	root
3	5	the	det
4	5	video	nn
5	2	lecture	obj


Parsed (i.e. predicted)

to	from	word	type
1	2	She	nsubj
2	0	saw	root
3	4	the	det
4	5	video	nsubj
5	2	lecture	ccomp

$$Acc = \frac{\#true\ dependencies}{\#dependencies}$$

$$UAS = \frac{4}{5} = 80\%$$

Evaluation – LAS



Labeled Attachment Score (LAS)

Gold (i.e. ground truth)

to	from	word	type
1	2	She	nsubj
2	0	saw	root
3	5	the	det
4	5	video	nn
5	2	lecture	obj

Parsed (i.e. predicted)

to	from	word	type
1	2	She	nsubj
2	0	saw	root
3	4	the	det
4	5	video	nsubj
5	2	lecture	ccomp

$$Acc = \frac{\#true\ dependencies}{\#dependencies}$$

$$UAS = \frac{4}{5} = 80\%$$

$$LAS = \frac{2}{5} = 40\%$$



State of the art

Bigger, deeper networks &
better hyperparameter tuning.

Beam search.

Self-attention.

Penn Treebank

Models are evaluated on the [Stanford Dependency](#) conversion (v3.3.0) of the Penn Treebank with **predicted** POS-tags. Punctuation symbols are excluded from the evaluation. Evaluation metrics are unlabeled attachment score (UAS) and labeled attachment score (LAS). UAS does not consider the semantic relation (e.g. Subj) used to label the attachment between the head and the child, while LAS requires a semantic correct label for each attachment. Here, we also mention the predicted POS tagging accuracy.

Model	POS	UAS	LAS	Paper / Source	Code
Label Attention Layer + HPSG + XLNet (Mrini et al., 2019)	97.3	97.42	96.26	Rethinking Self-Attention: Towards Interpretability for Neural Parsing	Official
ACE + fine-tune (Wang et al., 2020)	-	97.20	95.80	Automated Concatenation of Embeddings for Structured Prediction	Official
HPSG Parser (Joint) + XLNet (Zhou et al, 2020)	97.3	97.20	95.72	Head-Driven Phrase Structure Grammar Parsing on Penn Treebank	Official
Second-Order MFVI + BERT (Wang et al., 2020)	-	96.91	95.34	Second-Order Neural Dependency Parsing with Message Passing and End-to-End Training	Official
CVT + Multi-Task (Clark et al., 2018)	97.74	96.61	95.02	Semi-Supervised Sequence Modeling with Cross-View Training	Official
CRF Parser (Zhang et al., 2020)	-	96.14	94.49	Efficient Second-Order TreeCRF for Neural Dependency Parsing	Official
Second-Order MFVI (Wang et al., 2020)	-	96.12	94.47	Second-Order Neural Dependency Parsing with Message Passing and End-to-End Training	Official
Left-to-Right Pointer Network (Fernández-González and Gómez-Rodríguez, 2019)	97.3	96.04	94.43	Left-to-Right Dependency Parsing with Pointer Networks	Official

Image credit:

https://github.com/sebastianruder/NLP-progress/blob/master/english/dependency_parsing.md

More linguistic Features



Enriched features

Raw text presents challenges for NLP: rare words, varying word forms with similar meanings, & order-dependent interpretations. Enriching text with linguistic features can improve BOW models. This involves incorporating information about word function & relationships, such as POS tags & dependency relations, to enhance the model.



Word shapes

Problem

- Many words are rare.

Solution

- Featurize words by mapping them to simpler representations that captures word-attributes like length, capitalization, numerals, internal punctuation, Greek letters, etc.

Word	shape
Varicella-zoster	Xx - xxx
mRNA	xXXX
CPA1	XXXd



Word shapes

Featurize words by mapping them to simpler representations that captures word-attributes like length, capitalization, numerals, internal punctuation, Greek letters, etc.

First & last 2 characters are most important. The middle of words are replaced with unique codes in a specific order.

A, B, C, ..., Z \rightarrow x
a, b, c, ..., z \rightarrow x
0, 1, 2, ... \rightarrow d
-, :, ., ... \rightarrow -, :, ., ...

Word	shape
Varicella-zoster	Xx-xxx
mRNA	xXXX
CPA1	XXXd

Word substrings

Examples

- oxa – great precision for **drug** names, e.g. *Cotrimoxazole*.
- : – good precision for **movie** names, e.g. *Alien Fury: Countdown to Invasion*.
- field – ok precision for **person** names, e.g. *Wethersfield* & *Mansfield* & *Hadfield*.





Tutorial

MMAI5400_class06_enrichFeats.ipynb